
Question 1

(a) Is it possible to encapsulate an AH packet within an ESP packet and vice versa? Explain with the help of examples.

Consider the scenario (similar to Figure 17-3 in the text) where Alice and Bob are using IPsec and two firewalls, F1 and F2, between them are also using IPsec.

First, the scenario where the firewalls are using AH packets and Alice/Bob are using ESP packets, both using tunnel mode.

| | | | | |
|---------------------------|----|-------------------------|-----|--|
| IP: src = F1, dst = F2 | AH | IP: src = A, dst = B | ESP | |
|---------------------------|----|-------------------------|-----|--|

This scenario does **not** work. One difference of the integrity protection provided by AH/ESP is that AH provides integrity protection of the IP header. The IPsec between F1 and F2 cannot perform the integrity protection of Alice and Bob's IP header as it is inside an encrypted ESP packet and the F1/F2 do not have the key.

To get this scenario to work, Alice and Bob need to use transport mode. In transport mode the original IP header is not encrypted. This allows AH to perform its integrity protection of the IP header.

Alternatively, the firewalls could be using ESP packets and Alice/Bob are using AH packets.

| | | | | |
|---------------------------|-----|-------------------------|----|--|
| IP: src = F1, dst = F2 | ESP | IP: src = A, dst = B | AH | |
|---------------------------|-----|-------------------------|----|--|

In this scenario, the encapsulation of an AH packet inside a ESP packet **will** work. The IPsec between the routers are using ESP and do not check the IP headers during the integrity process.

(b) Suppose Alice is sending packets to Bob using IPsec. Suppose Bob's TCP acknowledgement gets lost, and Alice's TCP, assuming the packet was lost, retransmits the packet. Will Bob's IPsec implementation notice that the packet is a duplicate and discard it?

Bob's IPsec implementation will **not** notice it as a duplicate packet. Both AH and ESP IPsec headers have a *sequence number* field that function identically. This sequence number is used by IPsec to prevent replay attacks.

TCP handles the retransmission of packets. When TCP retransmits a packet, AH/ESP will treat it as a new packet with a new sequence number. This is one of the pros of building IPsec on top of TCP: it can hand-off the responsibility of guaranteed delivery (and discarding duplicate packets).

(c) Referring to Figure 17-2, assume that A and B are using IPsec in transport mode, and F1 and F2 have established an encrypted tunnel using IPsec. Assume A sends a TCP packet to B. Show the relevant fields of the IP header(s) as given to A's IPsec layer, as transmitted by A, as transmitted by F1, and as received by B.

Assume that both A/B and F1/F2 are using ESP packets.

Given to A's IPsec layer:

| | |
|---------------------|---|
| source address | A |
| destination address | B |

Transmitted by A:

Transmitted by F1:

Received by B:

(d) Why is IPsec not firewall-friendly?

A firewall looks at the fields in the transport layer (such as port and protocol) as part of its filtering process. Using the ESP protocol with IPsec encrypts this information and the firewall does not have the ability to decrypt this information. Solutions include terminating IPsec at the firewall or trusting the endpoint and forwarding the encrypted packet to the endpoint.

Question 2

(a) Why does the HMAC computation in SSL data transfer phase use a sequence number with each record even though SSL is built on top of TCP that delivers data in the correct order?

(b) This question pertains to 3G-UMTS security. Consider the case when a user visits another country and turns her phone on to make a phone call and notices service X. However, it is possible that service provider X is actually service provider Y that charges a lot more for phone calls made while roaming. The service provider Y when working as the intermediary between the subscriber and the user's home network appears to be X to the user by stays as Y for the home network. How does the home network ensure that the user knows the service provider it is actually connecting to?

(c) What are two advantages of computing the response to the challenge in the SIM card for GSM authentication?

(d) What are two problems in deploying DDoS prevention solutions close to the sources of attack traffic?

(e) Could the Local Aggregate Congestion Control mechanism for DDoS prevention result in collateral damage? Explain your answer.

(f) Would it make sense to use crypto cookies in conjunction with IP traceback? Explain briefly.

(g) What is a bloom filter? How could it be used to filter IP packets?

(h) Why doesn't the Botz-for-sale system use only one TCP connection to send the puzzle and receive the response instead of using two TCP connections?

(i) Why does it make sense to use an onion that encrypts information using public keys in the connection set up phase but secret keys during data movement?

Question 3

(1) What are the advantages of building security at the TCP layer?

There are many advantages to building security at the TCP layer. Since TCPsec is built on top of TCP, it can take advantage of some of the features of TCP. TCPsec can use TCP's guarantee of in-order delivery. Additionally it can use TCP's handshake phase to exchange keys used for TCPsec, removing an extra handshake phase.

By moving security to the transport layer, it provides end-to-end security. Maybe applications already use TCP and can update their kernel and set the TCP flags for quick deployment.

(2) What flexibility does TCPsec provide for encrypting the TCP segment?

Similar to SSL, TCPsec provides flexibility to the user for how to encrypt the TCP segment. First, encryption and integrity protection are independent operations and can be done in any order. Second, the user can specify which crypto algorithms they would like to be used for these operations. As for the actual encryption, the TCP payload is always encrypted. TCPsec leaves the option up to the user to encrypt portions of the IP header as well.

(3) How does TCPsec interoperate with NATS?

During the TCPsec handshake, a NAT option is provided. To detect if there are NATS (aka there are NATS on the path to the destination) the first SYN message is used. The NAT option contains the source IP address and port numbers. TCPsec can detect if these are changed (by a NAT) along the path to the destination. If detected, these new values can be used during message digest integrity checks and TCPsec plays well with NATS.

(4) Why does TCPsec handshake perform better than SSL?

SSL is deployed at the application layer. To perform an SSL handshake, a TCP connection is opened and the SSL handshake is performed over this connection. This TCP connection needs to perform the typical 3-way handshake to establish the connection. Once established, the application can use this connection to perform the SSL handshake, which requires multiple TCP packets to be transmitted back and forth.

In TCPsec, the security handshake is performed at the same time as the TCPsec 4-way handshake (a modified version of the 3-way TCP handshake). As you could imagine, this will be more performant than the SSL handshake.